When the Tenant Becomes the Landlord

Instructor Framing Guide | Discussion Questions Companion

Private teaching note. This companion is intended for instructors using the public case in graduate strategy, platform economics, technology management, or executive education settings. It preserves the discussion questions and adds facilitation guidance that is intentionally excluded from the public web version.

The case is bounded by what has shipped between October 2025 and April 2026; the questions below are bounded by what reasonable readers can defensibly argue from that record. None of the questions has a single correct answer. Each opens analytical territory where the case has documented evidence but has stopped short of committing to a verdict, and each invites students to commit where the case has not.

The questions are organized into three clusters. The first cluster engages the architectural bets that distinguish the major configuration-layer participants and the structural dynamics that shape which bet pays off on what timeline. The second cluster engages the strategic geometry of incumbent response, with comparisons drawn from prior platform migrations where the analogies serve and where they break down. The third cluster engages the falsifiability of the case's central claims and forces the reader to make a practitioner decision rather than rest in the analytical posture.

Cluster A: Architectural Bets and Configuration Dynamics

1. Anthropic has bet on the protocol-and-portfolio architecture: Model Context Protocol as cross-vendor agent-integration standard, multi-cloud distribution through Bedrock and Vertex AI and Foundry, and a portfolio of agentic execution products (Claude Code, Cowork, the Chrome browser agent). Google has bet on the internal-incumbent architecture: cross-application orchestration captured natively inside Workspace, advertising-revenue subsidy underwriting AIO deployment, and Android-layer agency rollout across Galaxy S26, Pixel 10, Android Auto, and Googlebook. OpenAI has bet on the superapp consolidation architecture: ChatGPT plus Codex plus Atlas integrated under a unified desktop platform, with consumer ChatGPT scale providing the conversion funnel. Which architectural bet is most defensible over a five-year horizon, and what evidence would falsify your answer within the first eighteen months?

INSTRUCTOR FRAMING

> The question forces students to commit on architectural bets the case has documented but not adjudicated. The defensibility arguments cut differently for each: the protocol-and-portfolio bet is harder to describe in revenue terms but easier to defend architecturally because it captures cross-vendor context-protocol rent regardless of which underlying model the user invokes; the internal-incumbent bet is harder to dislodge once installed but bounded by the Workspace customer base size relative to Microsoft's commercial base; the superapp consolidation bet captures the consumer-distribution funnel but concentrates execution risk on a single platform thesis. Productive disagreement among reasonable readers lives in how much weight to assign architectural defensibility versus revenue-line stability versus execution risk; assigning all three the same weight produces analytical mush and avoiding the trade-off is its own failure mode.

2. The developer-tooling slice has been in AIO execution longer and at greater depth than the office-productivity slice. The AIP-to-AIO transition that office-productivity has been crossing in 2026 happened in developer-tooling in 2024 and 2025, and AIO has been the dominant pattern there for over a year. Does the developer-tooling slice predict the office-productivity slice's trajectory, or do structural differences between the two slices break the analogy?

INSTRUCTOR FRAMING

> The case treats the slices as separate units of analysis where evidence supports it. Students should be able to name the structural differences (developer tooling lacks an entrenched paid-add-on revenue line analogous to the $30 Copilot seat; developers as a population are more tolerant of tool-switching than office workers; coding work has cleaner success criteria than knowledge-work synthesis) and weigh them against the structural similarities (both face the AIP-to-AIO transition, both face freemium-funnel pressure from entrants, both face the cost-to-value migration). The strongest answers will argue that the developer-tooling slice is a leading indicator on some dynamics and not on others, rather than treating it as universally predictive or universally distinct.

3. The freemium-funnel asymmetry Section 6.4 identifies is commercial rather than technical: Microsoft cannot make Copilot free in Office without destroying the $30 add-on, so the office product enters every competitive contest from behind the paywall while the entrants enter from in front of it. Construct the long-run unit-economics argument for the paid-only model surviving the freemium-funnel pressure, and identify the specific market conditions under which that argument holds.

INSTRUCTOR FRAMING

The question asks students to construct the steelman for the paid-only position rather than to dismiss it. The strongest versions invoke enterprise procurement cycles where freemium-conversion-funnel scale matters less than vendor-relationship depth, regulated industries where pre-licensed software simplifies compliance, integrated commercial relationships where the AI seat is a small component of a larger Microsoft license, and segments where the productivity gain from AI does not justify cross-vendor evaluation. The weak versions reach for "incumbents always win" without specifying the market conditions; the answer should name where the paid-only model survives and where it loses.

4. The Anthropic external-overlay AIO approach (cross-application context built over Microsoft's host applications without owning them) and Google's internal-incumbent AIO approach (cross-application context running as a native semantic layer over applications Google already owns) reach the same configuration through opposite directions. Read the parallel against the platform-research literature: which approach captures more durable platform position over the next decade, and how does the answer change if the underlying model layer commoditizes faster than expected?

INSTRUCTOR FRAMING

The question integrates Section 3.2's two-front advance with Section 4's Gawer-Cusumano frame and Section 8's capability-commoditization defeat condition. The external-overlay approach can ship AIO wherever the productivity suite is used, including environments Google and Microsoft do not control, but the cross-application context is parasitic on host applications the host vendor can change. The internal-incumbent approach ships natively with full integration quality and full control over host applications, but is locked to the host vendor's customer base. The commoditization branch matters because the internal-incumbent bet depends on the model layer remaining commercially differentiated; if the model layer commoditizes, the bundled pricing erodes and the external-overlay approach gains ground. Strong answers will distinguish the static comparison from the commoditization-conditional comparison.

Cluster B: Strategic Geometry and Incumbent Response

5. Section 4 introduces the absorbing-incumbent question because Google's structural position differs from the entrant positions Gawer-Cusumano predicts will produce asymmetric advantage through value-chain inversion. Argue whether Google's internal-incumbent move at the orchestration layer fits Christensen's incumbent-failure mode, escapes it, or invents a third pattern the existing frameworks do not capture. Cite the specific structural conditions that drive your answer.

INSTRUCTOR FRAMING

The question forces students to engage the limits of the case's analytical apparatus. Google's position is structurally distinct from both the Microsoft defensive position (the office-productivity revenue line foreclosing cross-application orchestration retrofit) and the Anthropic entrant position (external-overlay AIO across host applications it does not own). The strongest answers will name the specific structural conditions that distinguish Google's position: integrated advertising-revenue subsidy underwriting orchestration deployment, internal-incumbent ownership of the host applications themselves, multi-product portfolio spanning all three configuration tiers, and 250-million-vehicle Android Auto installed base extending agency-tier reach into territory the office-productivity vendors do not occupy. The weak answers will force Google into one of the existing two patterns rather than acknowledging the third.

6. Consider two platform-migration analogies. The Netscape-Microsoft browser war ended with Microsoft winning the browser war and losing the platform war: Internet Explorer captured the browser market by the early 2000s, but the platform power Microsoft attempted to retain through tying migrated to the web layer Netscape had pioneered, eventually to mobile and cloud layers Microsoft did not control. The Edge marginal-product trajectory: Microsoft continues investing in Edge despite sustained third-place global market share since 2015, because bundling generates enough revenue to justify the engineering investment. Argue whether Microsoft 365 Copilot is on the Netscape trajectory (winning the AIP layer but losing the platform war), on the Edge trajectory (sustained marginal-product position generating enough revenue to justify continued investment without capturing the platform), or on a third trajectory neither analogy captures.

INSTRUCTOR FRAMING

The question asks students to commit to a Microsoft trajectory using two historical analogies that constrain the answer space. The Netscape trajectory implies the platform power Microsoft attempts to retain at the AIP layer migrates to the AIO layer that the entrants are pioneering; the Edge trajectory implies Microsoft maintains a sustainable marginal-product position at the AIP layer without capturing the orchestration layer. The case's trifurcation analysis suggests a third trajectory: Microsoft loses at the office-productivity AIP slice, wins at the developer-tooling slice on freemium terms, and dominates at the enterprise-AI infrastructure slice in ways that fund continued investment across all three. The strongest answers will argue which trajectory captures Microsoft's most likely outcome and identify the specific evidence (current and prospective) that would distinguish the trajectories within the next two years.

7. Section 7.1 names Adobe, Salesforce, SAP, Oracle, and EZLynx as incumbent productivity vendors exposed to the same defensive-playbook pitfalls the Microsoft analysis surfaces. Select two of these vendors and argue whether the Microsoft canonical analysis (aggressive infrastructure spending, architectural diversification across configuration tiers) generalizes to each vendor's specific structural conditions, or whether the canonical analysis breaks down at the boundary of each vendor's business model.

INSTRUCTOR FRAMING

The question forces students to apply the trifurcation analysis to vendors with different structural conditions. Adobe has the creative-suite revenue line but lacks Microsoft's infrastructure dominance. Salesforce has the CRM lock-in but operates without the office-productivity revenue line. SAP has the enterprise-software relationship depth but lacks the consumer-facing AI surface. Oracle has the database installed base but operates inside business-critical workflows that resist disruption. EZLynx is the specialized-domain incumbent operating in the insurance agency vertical where domain depth substitutes for platform breadth. The strongest answers will identify which structural conditions each vendor can replicate from the Microsoft canonical analysis and which it cannot, and conclude with a defensible prediction for each selected vendor's configuration-migration outcome.

8. The configuration framework treats AIP, AIO, and AIA as categorical boundaries defined by structural conditions of membership rather than by enumeration of shipping examples. The open-source ecosystem ships capabilities (LangChain orchestration, AutoGPT-derived agentic systems, OpenInterpreter local execution) that satisfy the structural conditions for AIO and AIA without operating inside any vendor's commercial subscription. How does open-source AIA presence affect the rent-capture dynamics the case documents, and what conditions would have to hold for open-source AIA to constrain commercial-vendor pricing power at the configuration layer?

INSTRUCTOR FRAMING

The question opens territory the case acknowledges but does not engage: open-source presence at the configuration layer creates a pricing-power constraint that the commercial-vendor analysis does not fully internalize. The case argues that vendors price against value rather than against compute cost because the configuration position captures rent the user accepts in exchange for the dependency. Open-source alternatives that satisfy the structural conditions of AIO or AIA create a substitution threat at the high end of the value-anchored pricing tiers. The strongest answers will identify the specific conditions under which open-source AIA constrains commercial pricing (production-readiness parity, sufficient integration depth to overcome adoption friction, governance models that enterprise procurement accepts) and will name the segments where commercial vendors retain pricing power despite open-source presence (regulated industries, enterprise integration depth, support-and-warranty obligations, brand-trust signals).

Cluster C: Falsifiability and Practitioner Decisions

9. Section 7.5 argues that regulation will trail the technology by a full model generation and that the vendors who win the configuration-migration contest will win on facts on the ground established before any regulatory architecture catches up. Construct the strongest counter-argument: identify a specific regulatory action plausible within the next twenty-four months that would materially constrain the configuration-migration trajectory the case documents, and assess the probability of that action occurring before the configuration positions consolidate.

INSTRUCTOR FRAMING

> The question forces students to argue against the case's regulatory-trailing thesis with specific regulatory actions, not generic appeals to "regulators might act." Candidates include EU AI Act compliance enforcement against general-purpose AI providers, DMA gatekeeper designation extended to AI orchestration layers, FTC action under merger-review or consumer-protection authority, state attorney general coordinated action on AI consumer harms, or copyright-litigation outcomes producing structural remedies. The strongest answers will name the specific action, identify the legal vehicle, assess the institutional and political conditions required for action, and predict the timeline against which the configuration positions consolidate. The weak answers will treat regulation as a generic force without specifying mechanism.

10. The case identifies capability commoditization as a defeat condition for the configuration framework: if frontier-model capability commoditizes faster than configuration positions cement, the rent at each layer erodes, and the internal-incumbent bets are the most exposed because the bundled pricing depends on the model layer being commercially differentiated. Construct the specific scenario under which capability commoditization defeats the configuration migration, including the model-layer dynamics (open-source frontier-parity, multiple commercially viable providers at parity, the specific capability thresholds that trigger commoditization), the timeline over which commoditization would have to occur to defeat the migration, and the cross-vendor strategic responses you would expect to observe as commoditization unfolds.

INSTRUCTOR FRAMING

> The question opens the most analytically substantive of the three defeat conditions. The strongest answers will distinguish commoditization at different capability thresholds: commoditization at GPT-4-level capability is already largely true and has not defeated the migration; commoditization at the current frontier (GPT-5.4, Claude Opus 4.7, Gemini 3) would erode some configuration positions but not all; commoditization at the next frontier (assuming such a frontier exists) would erode the value-anchored pricing the case identifies. The strongest answers will also identify which vendors are most exposed to commoditization (the internal-incumbent bet is most exposed; the protocol-and-portfolio bet is least exposed; the superapp consolidation bet is intermediate) and will name the cross-vendor strategic responses (pricing compression, vertical integration into hardware or data, regulatory-protection-seeking, talent acquisition) that would be visible early indicators of commoditization-driven dislocation.

11. You are advising one of the five stakeholder categories Section 7 addresses: incumbent productivity vendor, frontier-model entrant, enterprise customer, workforce or educator, or regulator. Choose one category. Make a specific recommendation for what your client should do in the next twelve months in light of the case's analysis, and identify the two pieces of evidence that, if observed within the same twelve-month window, would force you to retract the recommendation and revise the strategy.

INSTRUCTOR FRAMING

> The closing question forces students out of the analytical posture into the practitioner posture. The recommendation must be specific (a concrete action within a defined timeline, not a posture or an orientation), and the retraction conditions must be observable evidence rather than analytical disagreement. The strongest answers will commit to a recommendation that exposes the student to being wrong and will identify the evidence that would prove them wrong; the weak answers will hedge until the recommendation becomes content-free or will name evidence so unlikely or so distant in time that the retraction condition is purely theoretical. Reasonable disagreement among instructors is welcome on which stakeholder category produces the most analytically substantive answers; the question is designed to work for all five.